

ZICE 17 On-boarding



Luca Mazzone & Simon Scheidegger
Jan. 24th, 2017
ZICE 17 – University of Zürich

Outline

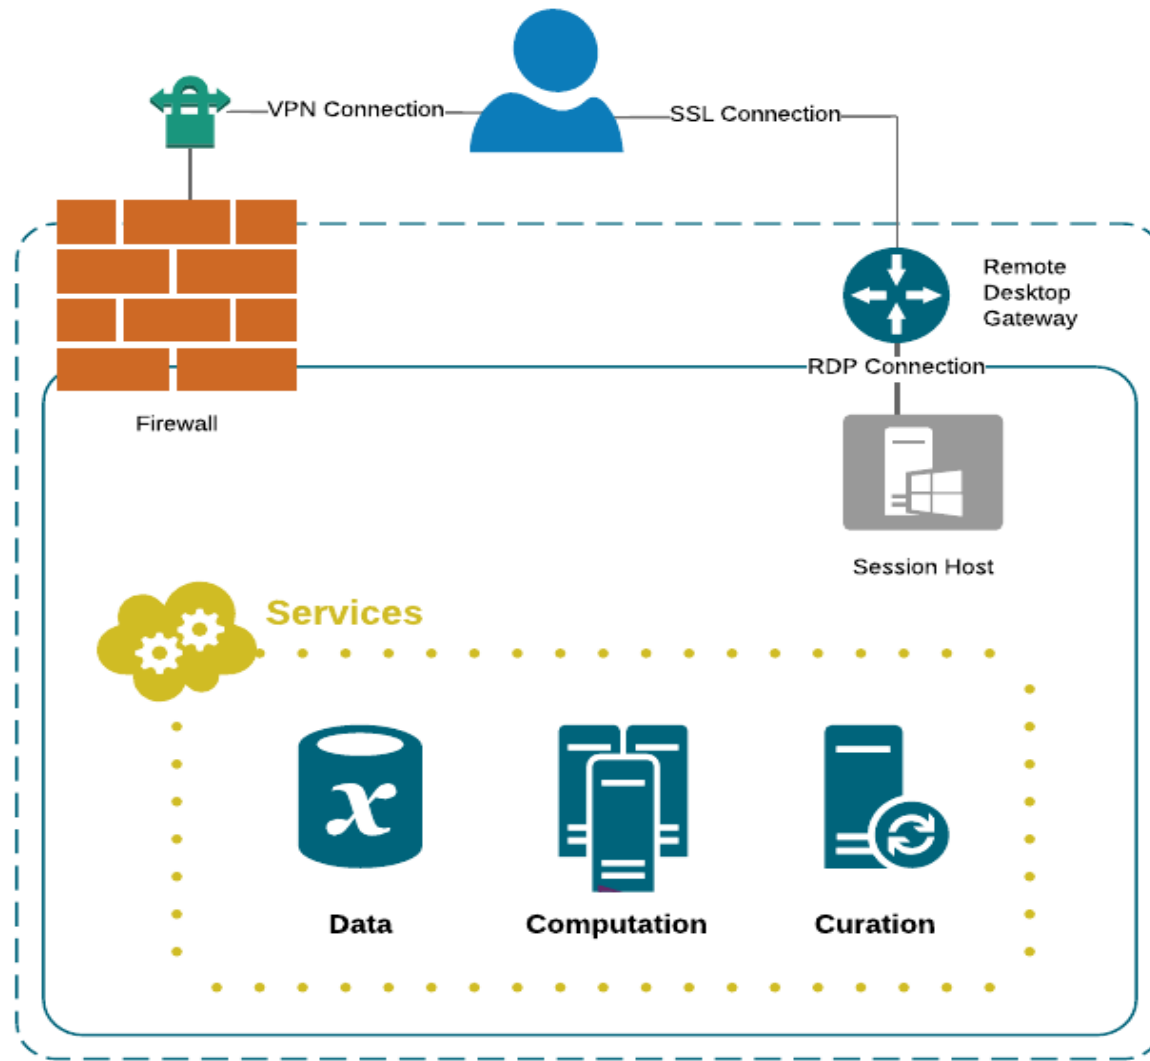
1. The compute infrastructure for ZICE
2. Accessing WLAN during ZICE
3. Access to ALPHACRUNCHER Services
4. Course management – git Classroom
5. First steps on a Linux Cluster

Aim


- Every student works on a unified environment (no issues with licenses etc.)
- Service to the community: knowledge transfer
Access to resources for an extended period of time, even outside the UZH network.

1. The compute infrastructure

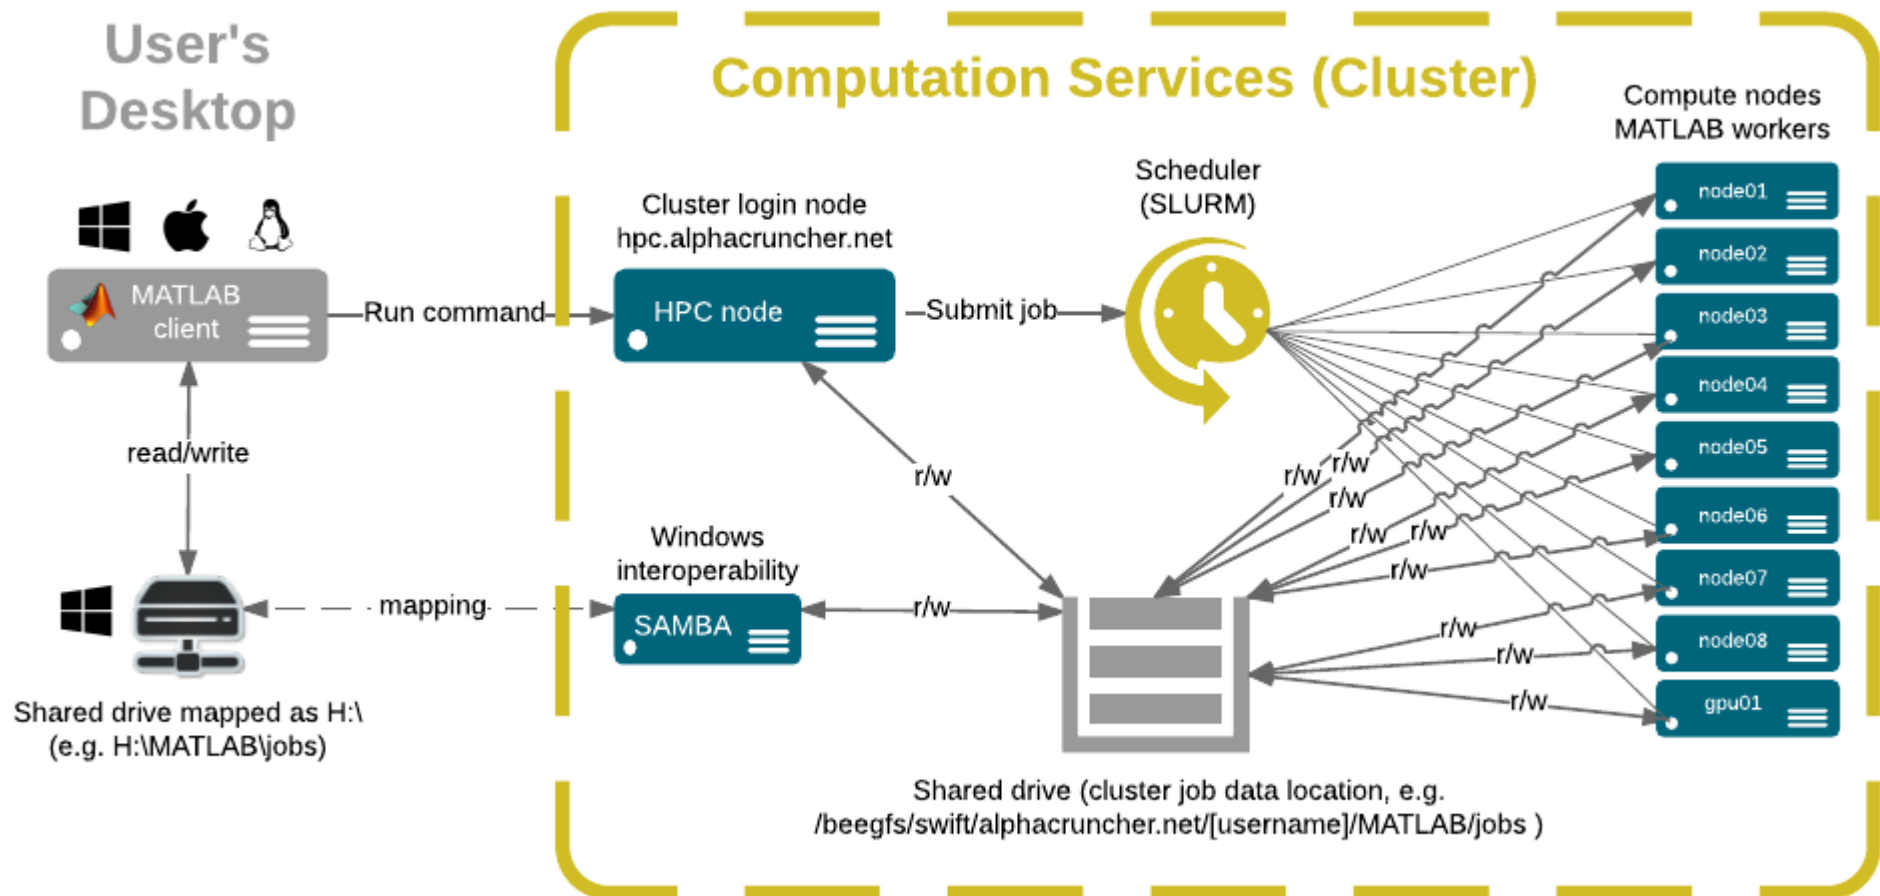
©Alphacruncher



Session-based desktop deployment



Compute Cluster



2. Accessing WLAN during ZICE

You all received the login credential today.

Make sure you **DO (IF POSSIBLE) NOT LOG INTO THE NETWORK WITH MULTIPLE DEVICES AT THE SAME TIME and no NETFLIX PLEASE :)**

→ we want to avoid bandwidth problems.

WiFi Names (SSID) – passwd: **IBF_router**

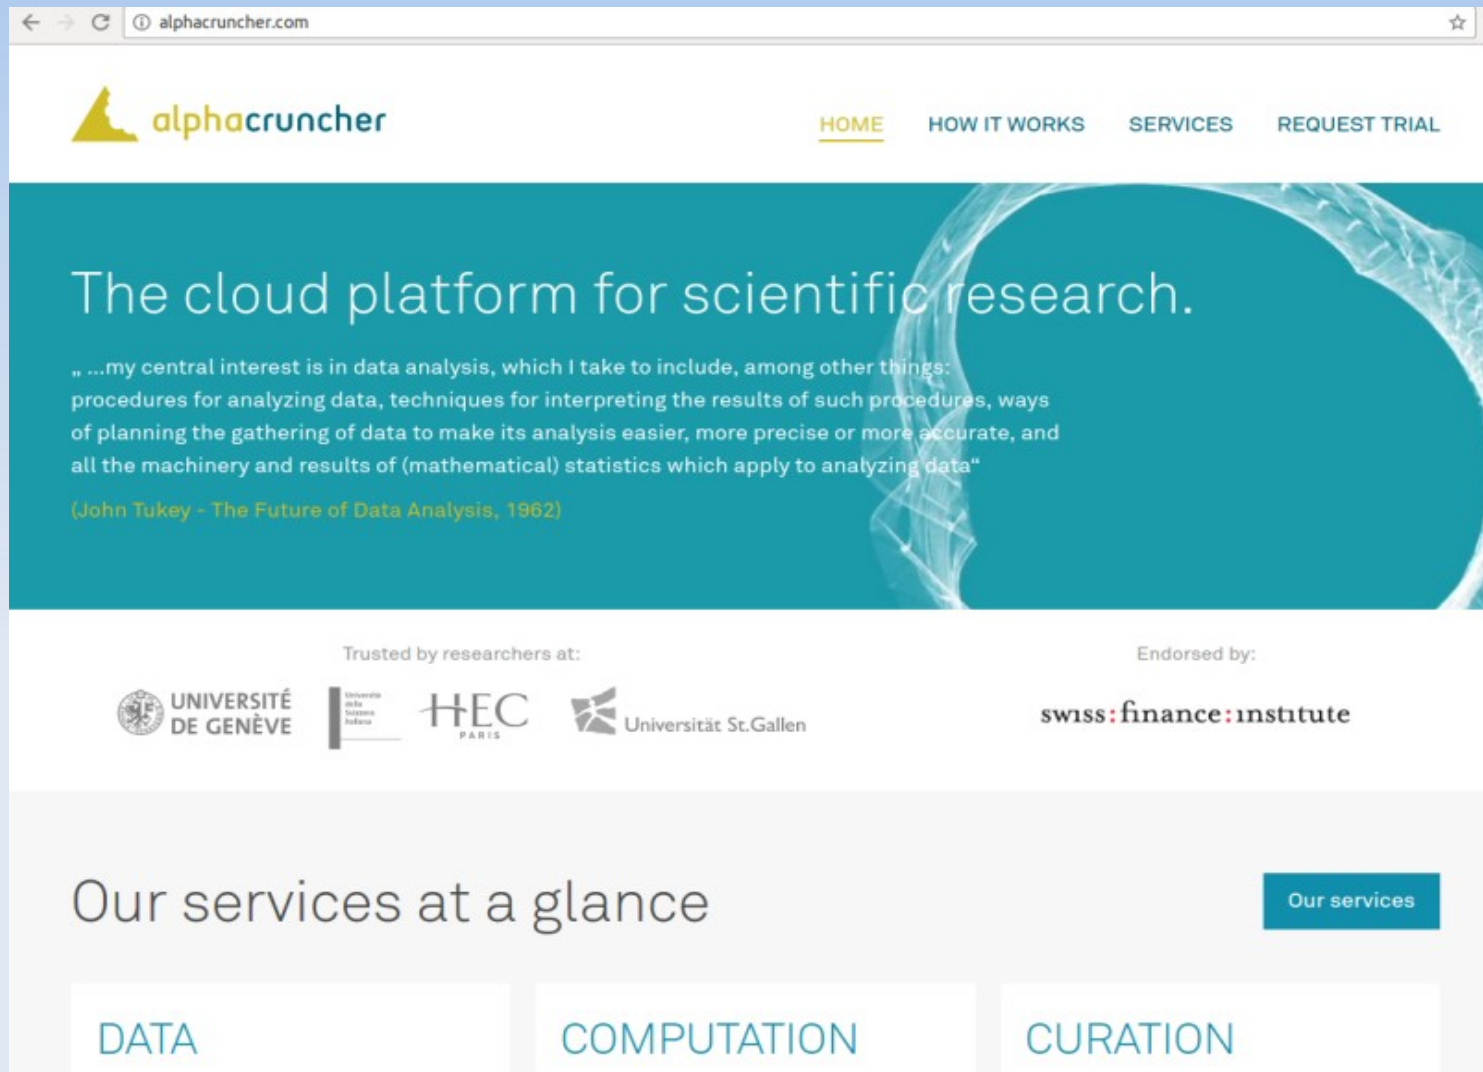
IBF Router 1

IBF Router 2

IBF Router 3


3. Access to Alphacruncher Services

URL: alphacruncher.com



The screenshot shows the homepage of alphacruncher.com. The browser address bar displays "alphacruncher.com". The website header includes the alphacruncher logo (a yellow triangle) and navigation links: HOME (underlined), HOW IT WORKS, SERVICES, and REQUEST TRIAL. The main content area has a teal background with a white line-art illustration of a human head profile. The text reads: "The cloud platform for scientific research." followed by a quote from John Tukey (1962) about data analysis. Below this, logos of partner institutions are shown under the headings "Trusted by researchers at:" and "Endorsed by:". The bottom section, titled "Our services at a glance", features three service categories: DATA, COMPUTATION, and CURATION, each in a white box with a teal border. A teal button labeled "Our services" is positioned to the right of the title.

← → ↻ ⓘ alphacruncher.com ☆

 **alphacruncher**




[HOME](#) [HOW IT WORKS](#) [SERVICES](#) [REQUEST TRIAL](#)

The cloud platform for scientific research.

„...my central interest is in data analysis, which I take to include, among other things: procedures for analyzing data, techniques for interpreting the results of such procedures, ways of planning the gathering of data to make its analysis easier, more precise or more accurate, and all the machinery and results of (mathematical) statistics which apply to analyzing data“

(John Tukey - The Future of Data Analysis, 1962)

Trusted by researchers at:

 **UNIVERSITÉ DE GENÈVE**  **HEC PARIS**  **Universität St.Gallen**

Endorsed by:

swiss:finance:institute

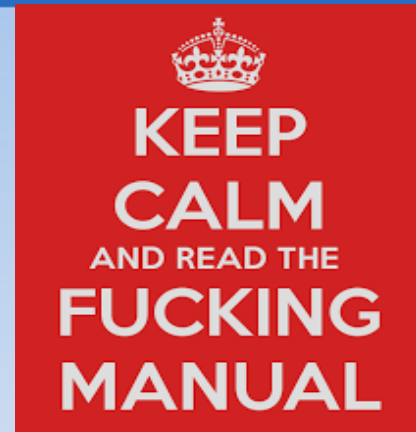
Our services at a glance

[Our services](#)

DATA **COMPUTATION** **CURATION**

Access to the remote desktop

- You all obtained an access guide for Windows – Mac OS X – Linux
- More in the lecture notes – GIT classroom



→ https://github.com/edualphacruncher/zice17-YOUR_GITBHUB_PRIVATE_REPOSITORY/simon/onboarding/onboarding_ZICE17.pdf

→ https://github.com/edualphacruncher/zice17-YOUR_GITBHUB_PRIVATE_REPOSITORY/zice17/simon/onboarding/access_guide_zice.pdf

- We will step you now through Windows & Mac OS X

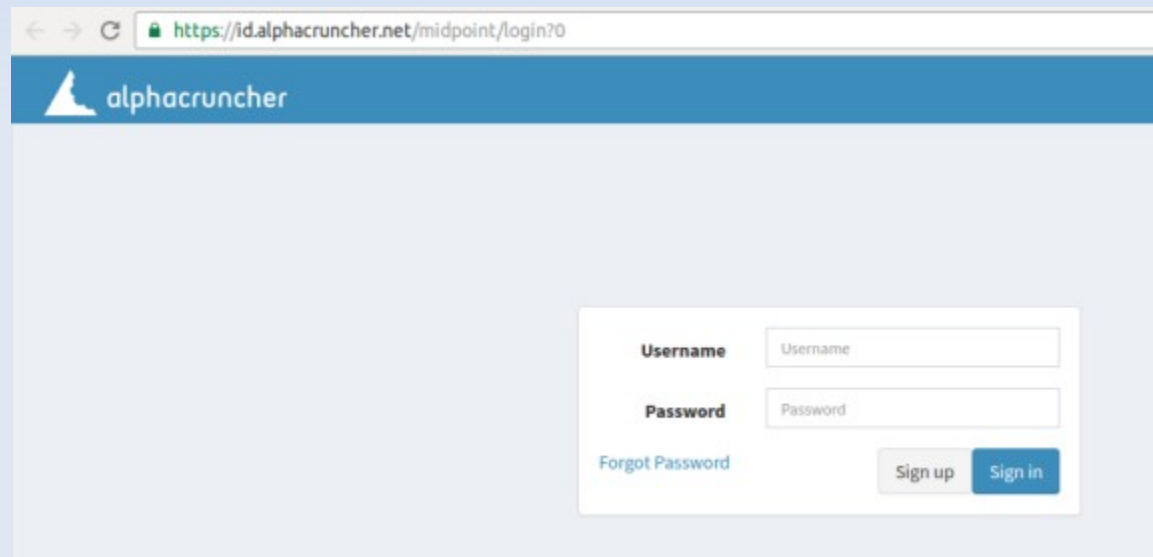
User name and password

You all (?) obtained an email stating

Access to the platform is available via Remote Desktop Connection. You need to (re)set the password first before proceeding. Please point your browser to <https://id.alphacruncher.net> and follow the instructions (see [access guide](#)).

Your username: **XYZ**

→ Reset your password here: **<https://id.alphacruncher.net>**
Click on the **Forgot Password** Link and follow the instructions



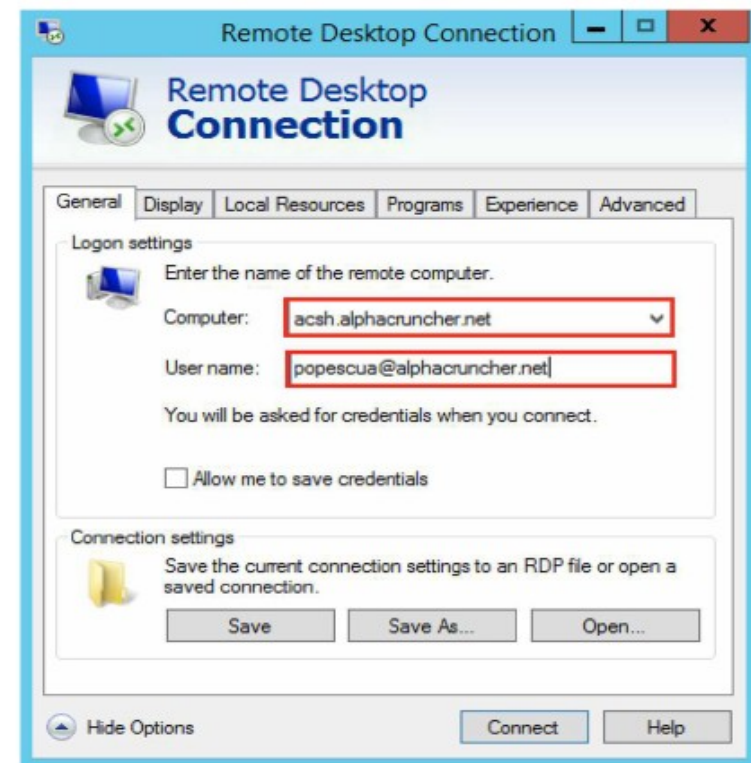
The screenshot shows a web browser window with the URL <https://id.alphacruncher.net/midpoint/login?0>. The page header features the 'alphacruncher' logo. The main content area contains a login form with the following elements:

- Username** label next to a text input field.
- Password** label next to a text input field.
- A **Forgot Password** link below the password field.
- Sign up** and **Sign in** buttons at the bottom right of the form.

Windows (I)

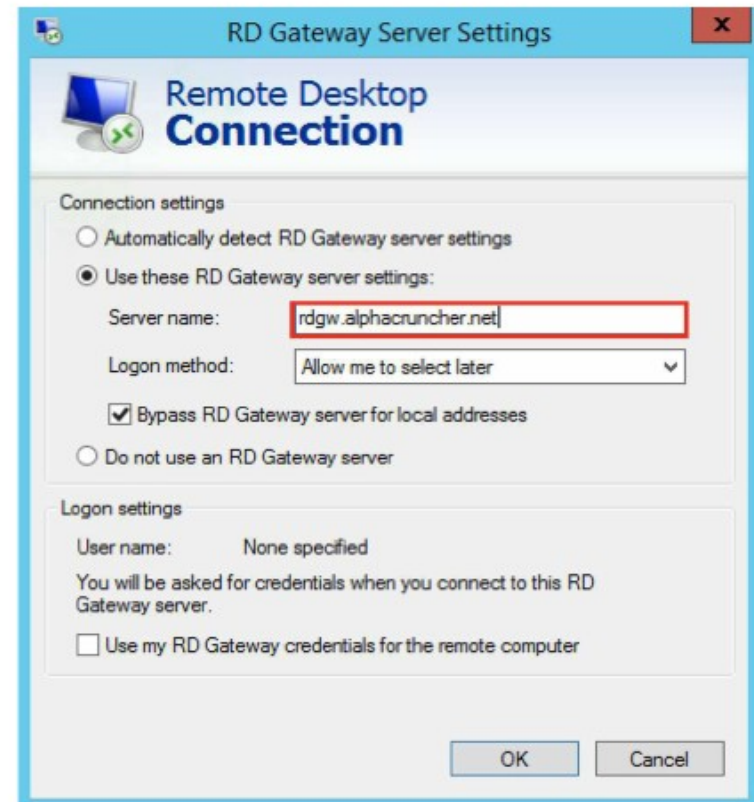
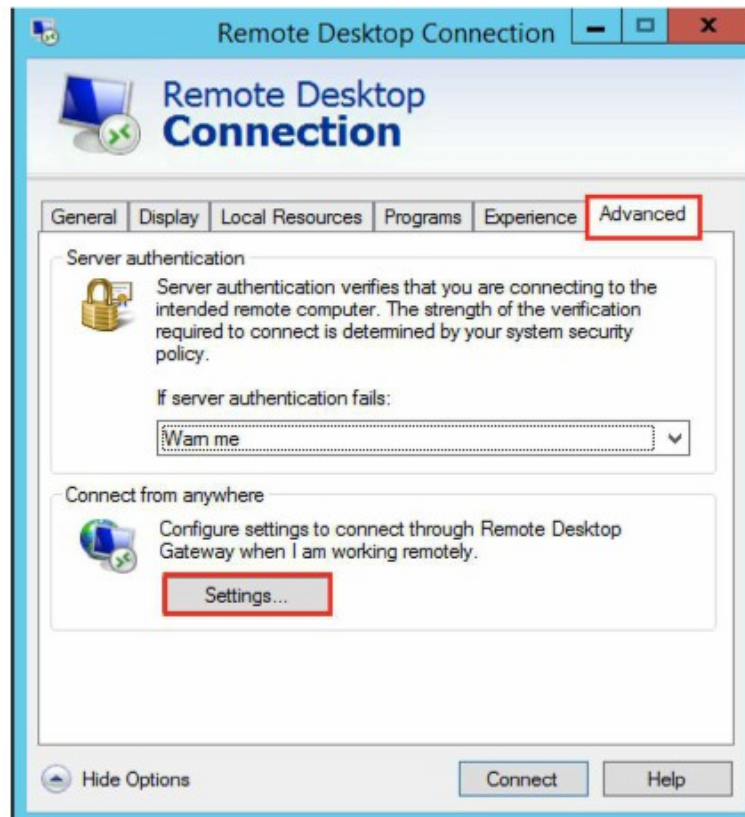


Access Remote Desktop Connection from the **Start** button, then click on **show options**



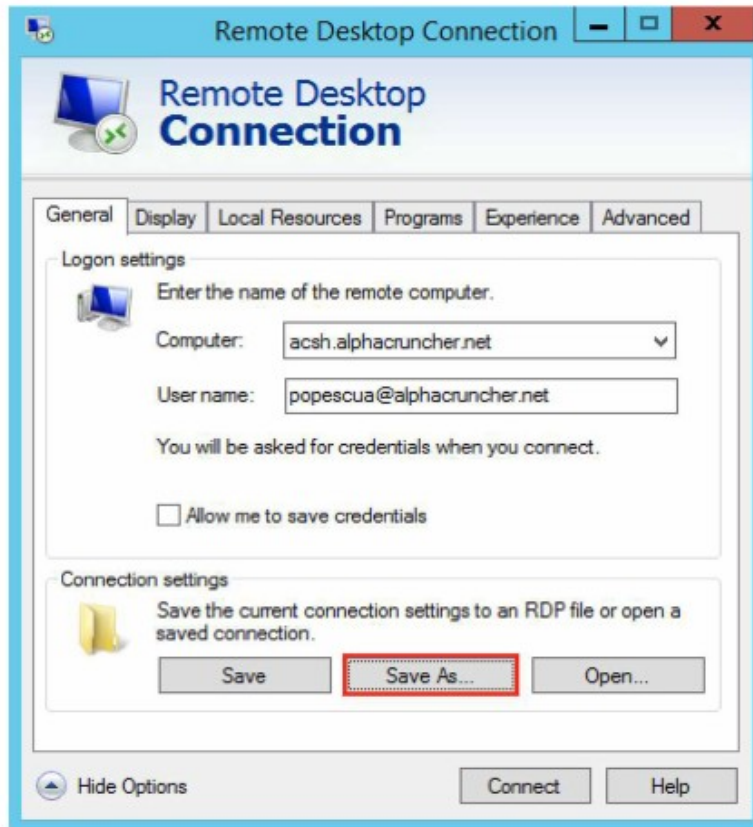
Computer: **eduzh.alphacruncher.net**
User Name: **your** username

Windows (II)




Go on "Advanced", click "Settings" Server: **rdgw.alphacruncher.net**
Logon: *Allow me to select later*

Windows (III)

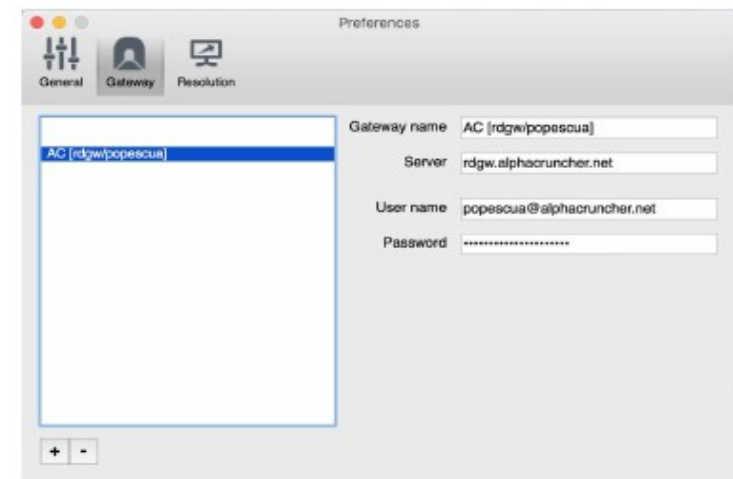



Click on "**save as**"



Save with ***.rdp** extension. Now you can connect

MAC OS X



Download the Remote Desktop Client from the Mac App Store.
Set up a New Gateway and click **Preferences**

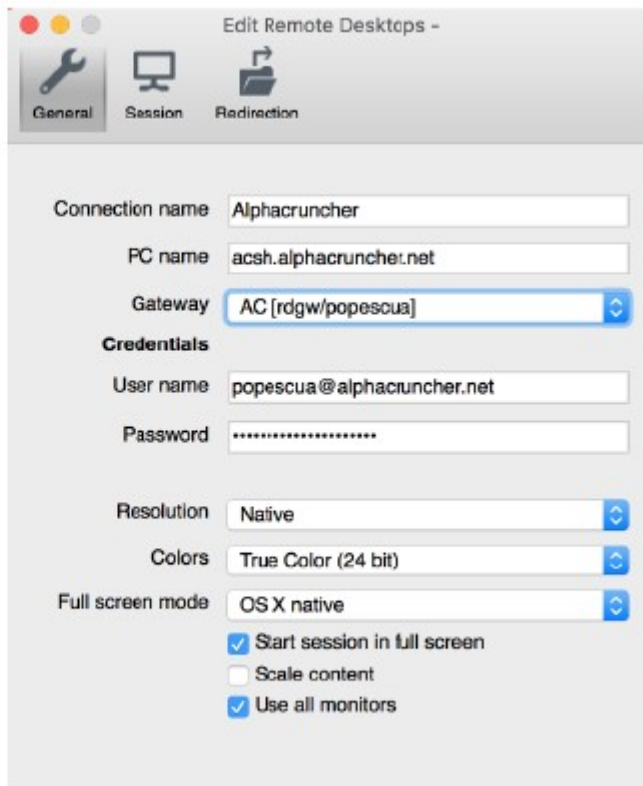
Gateway: AC[rdgw/username]
Server: rdgw.alphacruncher.net
User:
[username]@alphacruncher.net

MAC OS X (II)

On the Connection Center screen, click **New**

Enter the following information:


- General: Alphacruncher
 - PC name: eduzh.alphacruncher.net
 - Gateway: AC [rdgw/your username]
 - User name: [your username]@alphacruncher.net
- Connect!



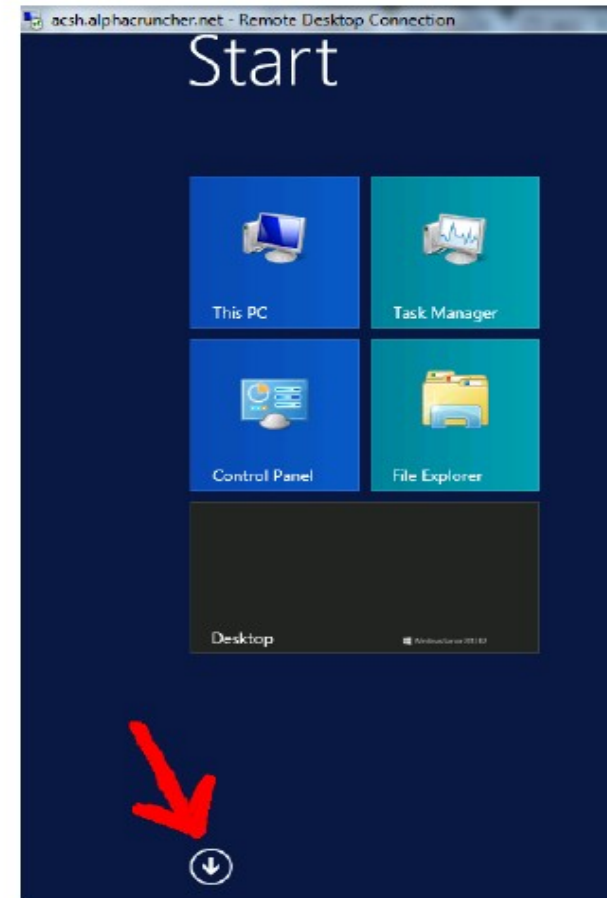
The screenshot shows the 'Edit Remote Desktops' window with the 'General' tab selected. The fields are as follows:

Field	Value
Connection name	Alphacruncher
PC name	acsh.alphacruncher.net
Gateway	AC [rdgw/popescua]
Credentials	
User name	popescua@alphacruncher.net
Password
Resolution	Native
Colors	True Color (24 bit)
Full screen mode	OS X native
<input checked="" type="checkbox"/> Start session in full screen	
<input type="checkbox"/> Scale content	
<input checked="" type="checkbox"/> Use all monitors	

Remote desktop

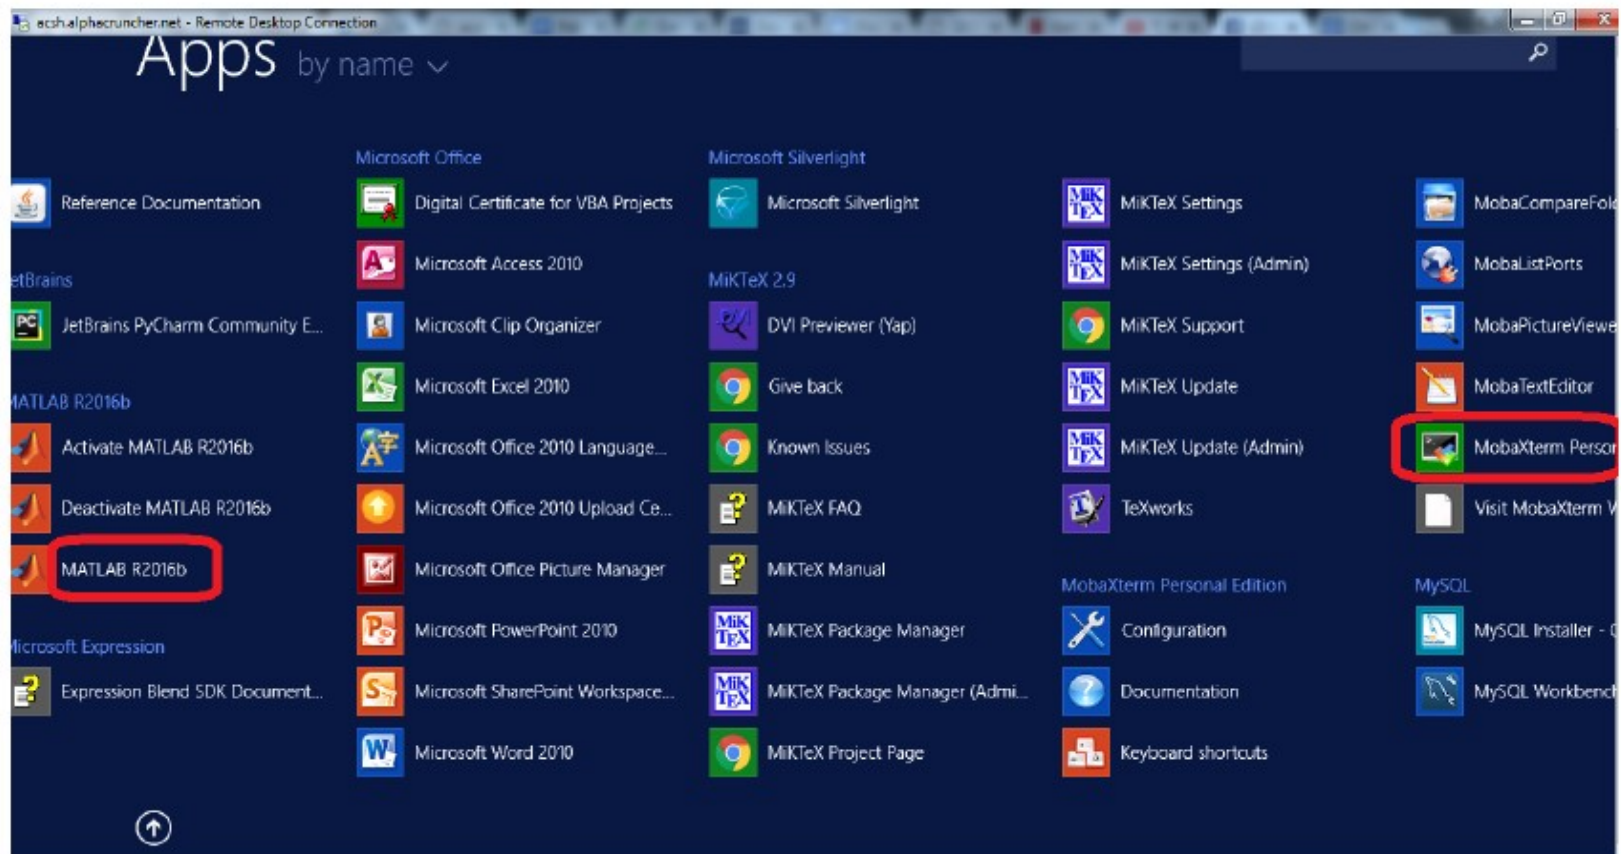


Click on the **Start** button



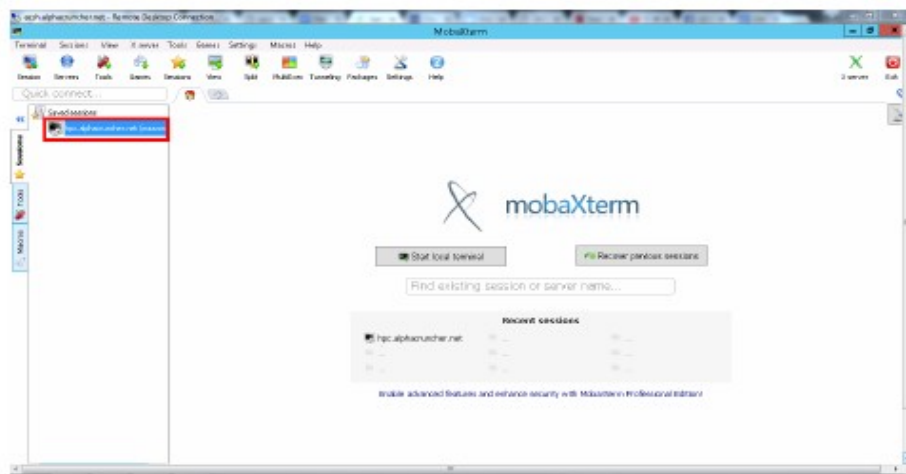
Use the lower-pointing arrow to scroll down...

Two apps you need all the time

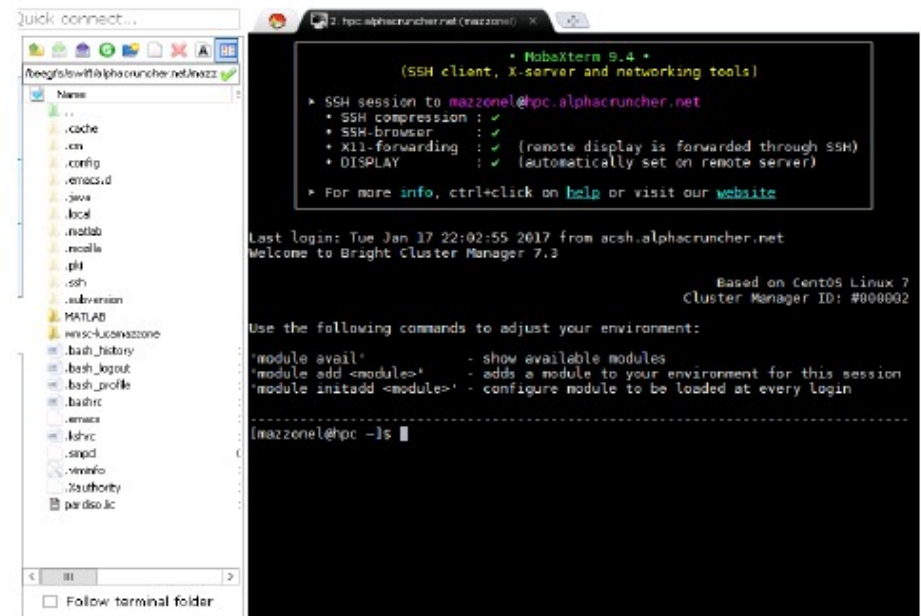


Now you can open Matlab 2016b in display mode by clicking on the icon
Click on MobaXterm to access the command line of the cluster

Accessing the “Linux” cluster



Access by clicking on the highlighted icon on top left...



... and you're finally on the command line of the cluster!

Help

On-sight support:

- Luca Mazzone (luca.mazzone@bf.uzh.ch)

Support for nasty cases:

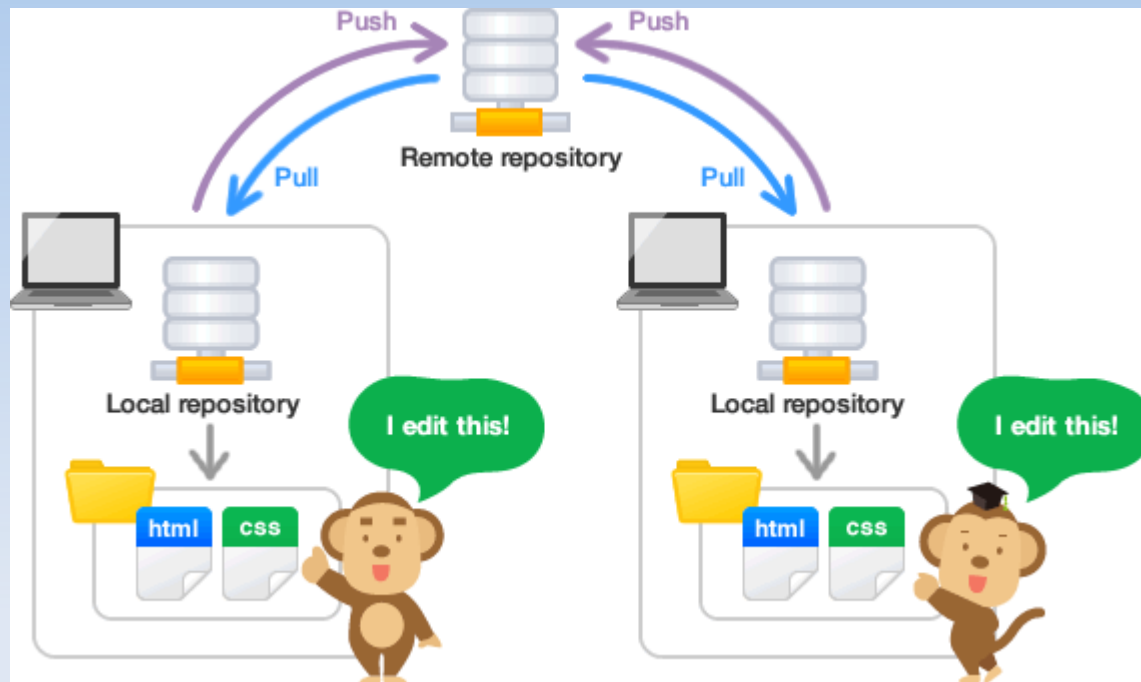
- Alphacruncher (support@alphacruncher.com)

Emergency support :)

<http://imgtfy.com/?q=help+on+alphacruncher>

4. Course management – GIT classroom

Link to classroom <https://github.com/edualphacruncher/zice17>



The official website: <https://git-scm.com/>

Free git book: <https://progit2.s3.amazonaws.com/en/2016-03-22-f3531/progit-en.1084.pdf>

How I use git :)




<https://xkcd.com/1597/>

Git overview

- As you develop software and make changes, add features, fix bugs, etc. it is often useful to have a mechanism to keep track of changes and to ensure that your code base and artifacts are well-protected by being stored on a reliable server (or multiple servers).
- This allows you access to historic versions of your application's code in case something breaks or to “roll-back” to a previous version if a critical bug is found.
- The solution is to use a revision control system that allows you to “check-in” changes to a code base.
- It keeps track of all changes and allows you to “branch” a code base into a separate copy so that you can develop features or enhancements in isolation of the main code base (often called the “trunk” in keeping with the tree metaphor).
- Once a branch is completed (and well-tested and reviewed), it can then be merged back into the main trunk and it becomes part of the project.

Version control system

- Git essentially keeps track of all changes made to a project and allows users to work in large teams on very complex projects while minimizing the conflicts between changes.
- These systems are not only used for organizational and backup purposes, but are absolutely essential when developing software as part of a team.
- Each team member can have their own working copy of the project code without interfering with other developer's copies or the main trunk.
- Only when separate branches have to be merged into the trunk do conflicting changes have to be addressed.
- Otherwise, such a system allows multiple developers to work on a very complex project in an organized manner.

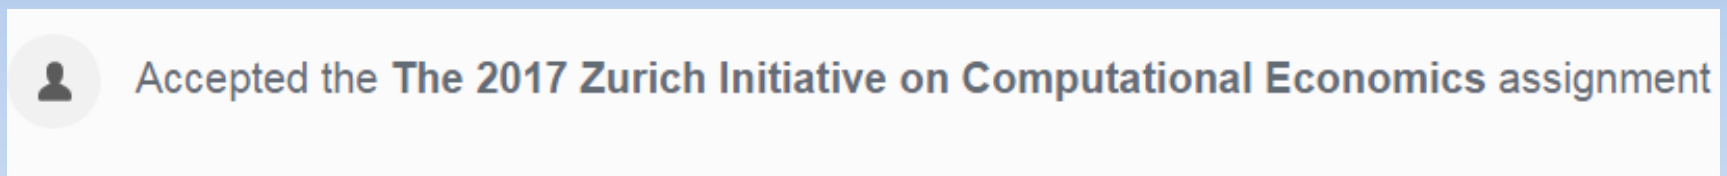


ZICE 17 Classroom

- We use git to provide you with **lecture slides/codes**
- GitHub assignment link for class:
<https://classroom.github.com/assignment-invitations/2dcd7145a43f8b7daba3e77fa84d88c>
- This is a fork of the master repository:
<https://github.com/edualphacruncher/zice17>
(Private repository available only to speakers)
- Keep the master repository up to date (cf. Slide 28)
- You find github's cheat sheet here:
<.../zice17/simon/onboarding/github-git-cheat-sheet.pdf>

Git on the remote desktop

You received an e-mail containing an invitation to participate in the GitHub repository. By clicking it, you will see that you have:



In the following e-mail, copy the address of your private repository in the clipboard – you'll need this later.

Hey there, we're just writing to let you know that you've been automatically subscribed to a repository on GitHub.

edualphacruncher/zice17-lucamazzone created by popescua

zice17-lucamazzone created by GitHub Classroom

<https://github.com/edualphacruncher/zice17-lucamazzone>



you need this address!

You'll receive notifications for all issues, pull requests, and comments that happen inside the repository. If you would

Notice: the address will be in the form

<https://github.com/edualphacruncher/zice17-YOUR-PRIVATE-REPOSITORY>

Git on the remote desktop (II)

On the first day, you need to create a corresponding folder where to receive files.


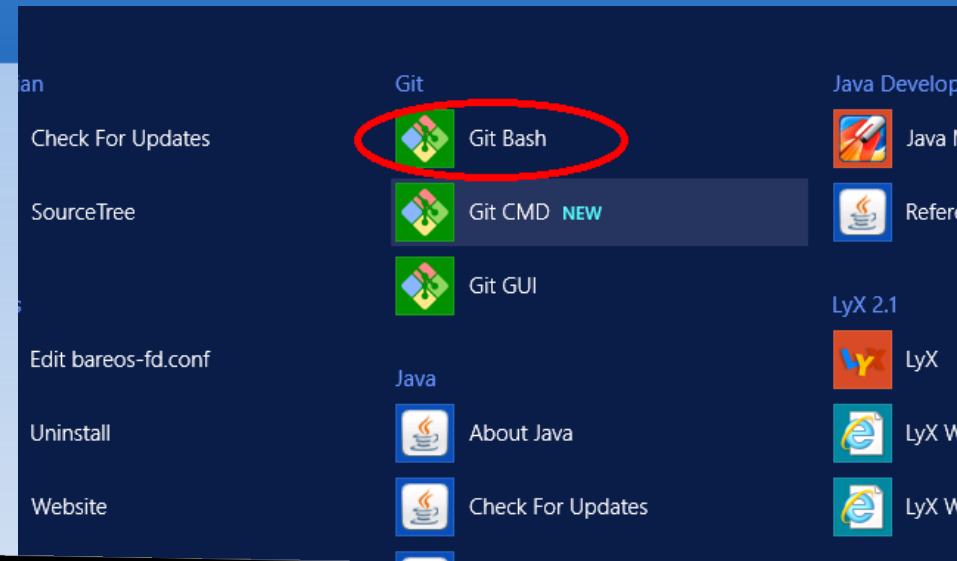
On the remote computer command line, write:

```
$ cd Desktop  
$ git clone address
```

Where the address is again:

<https://github.com/edualphacruncher/...>

[zice17-YOUR-PRIVATE-REPOSITORY](#)



Git on the remote desktop (III)

To keep your copy of the ZICE17 master repository up to date:

→ Change the current working directory to your local private repository:

```
mazzonel@eduzh MINGW64 ~/Desktop (master)
$ cd zice17/
```

→ List the current configured remote depository for your fork:

```
mazzonel@eduzh MINGW64 ~/Desktop/zice17 (master)
$ git remote -v
origin  git@github.com:edualphacruncher/zice17.git (fetch)
origin  git@github.com:edualphacruncher/zice17.git (push)
```

→ Specify a new remote upstream repository that will be synchronized with your fork:

```
mazzonel@eduzh MINGW64 ~/Desktop/zice17 (master)
$ git remote add upstream git@github.com:edualphacruncher/zice17
```

→ Verify the new upstream repository you have specified:

```
mazzonel@eduzh MINGW64 ~/Desktop/zice17 (master)
$ git remote -v
origin  git@github.com:edualphacruncher/zice17.git (fetch)
origin  git@github.com:edualphacruncher/zice17.git (push)
upstream      git@github.com:edualphacruncher/zice17 (fetch)
upstream      git@github.com:edualphacruncher/zice17 (push)
```

Git on the remote desktop (IV)

Once you synchronized the remote repository to your local folder, you can fetch from the remote repository whenever there is an update (this will happen throughout the workshop, and you will be instructed by lecturers about when to update your repository to get the new materials!). To fetch and update, do the following:

Fetch the branches and their respective commits from the upstream repository:

```
mazzonel@eduzh MINGW64 ~/Desktop/zice17 (master)
$ git fetch upstream
remote: Counting objects: 12, done.
remote: Compressing objects: 100% (10/10), done.
remote: Total 12 (delta 3), reused 3 (delta 0), pack-reused 0
Unpacking objects: 100% (12/12), done.
From github.com:edualphacruncher/zice17
* [new branch]      master      -> upstream/master
```

Finally, update your private repository!

```
mazzonel@eduzh MINGW64 ~/Desktop/zice17 (master)
$ git pull upstream master
From github.com:edualphacruncher/zice17
* branch            master      -> FETCH_HEAD
Updating c3b0066..d1b5ef1
Fast-forward
 ZICE17_info/README.md | 5 +++++
simon/onboarding/onboarding_ZICE17.pdf | Bin 1991822 -> 0 bytes
2 files changed, 5 insertions(+)
create mode 100644 ZICE17_info/README.md
delete mode 100644 simon/onboarding/onboarding_ZICE17.pdf
```

Git on the cluster

MobaXTerm, log in

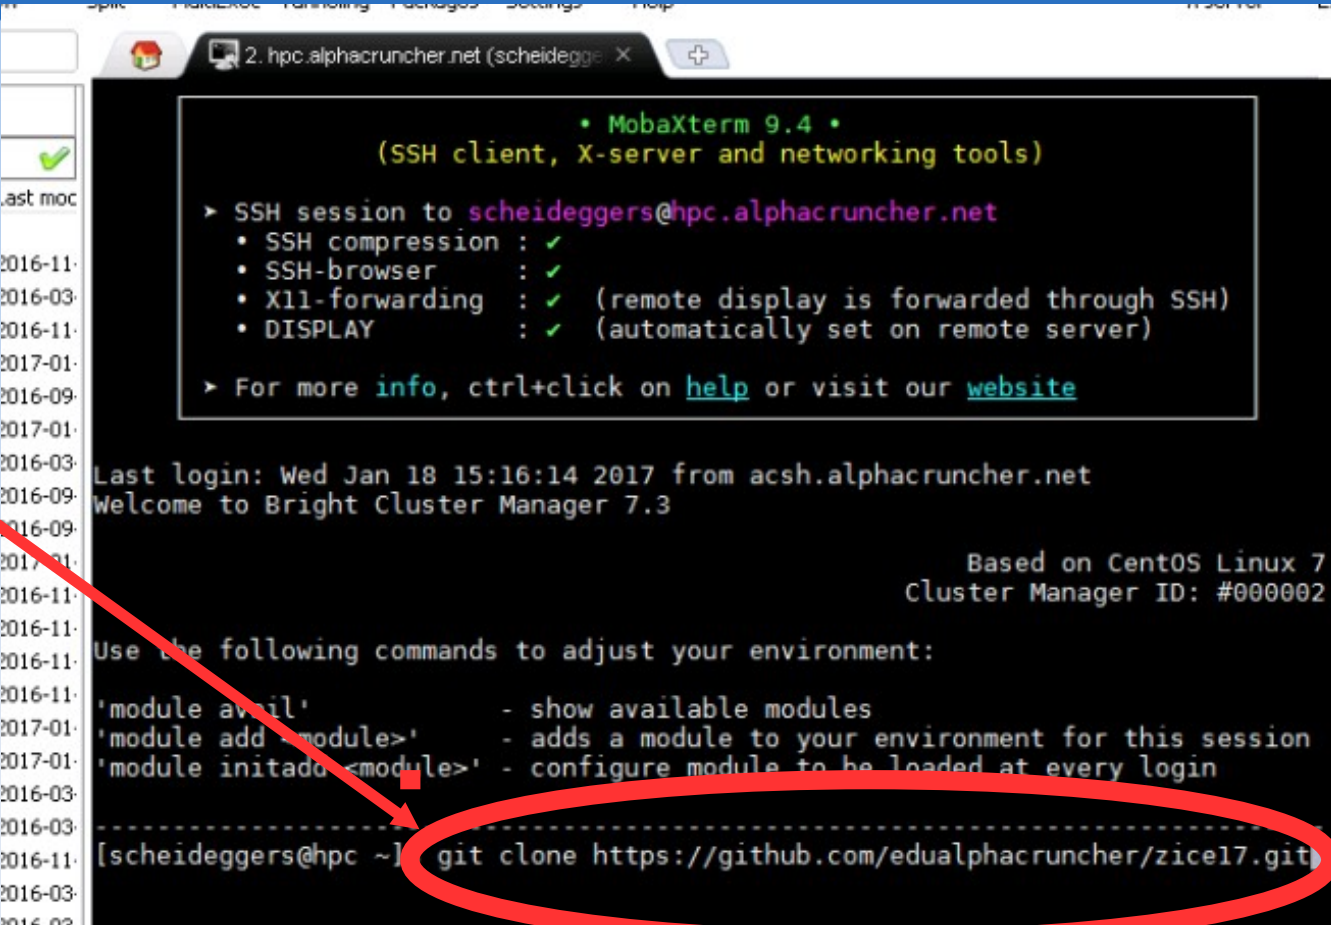
clone the repository:
\$ git clone *address*

To update the repository
in future circumstances,
go to the directory

\$ cd zice17-YOUR-PRIVATE-REPOSITORY

Then you can follow exactly the
same directions as in the
previous slides!

Again, *address* is: <https://github.com/edualphacruncher/zice17-YOUR-PRIVATE-REPOSITORY>



```
2. hpc.alphacruncher.net (scheidegger X)
• MobaXterm 9.4 •
(SSH client, X-server and networking tools)

> SSH session to scheideggers@hpc.alphacruncher.net
• SSH compression : ✓
• SSH-browser      : ✓
• X11-forwarding   : ✓ (remote display is forwarded through SSH)
• DISPLAY          : ✓ (automatically set on remote server)

> For more info, ctrl+click on help or visit our website

Last login: Wed Jan 18 15:16:14 2017 from acsh.alphacruncher.net
Welcome to Bright Cluster Manager 7.3

Based on CentOS Linux 7
Cluster Manager ID: #000002

Use the following commands to adjust your environment:

'module avail' - show available modules
'module add <module>' - adds a module to your environment for this session
'module initadd <module>' - configure module to be loaded at every login

[scheideggers@hpc ~]$ git clone https://github.com/edualphacruncher/zice17.git
```

4. First steps on a Linux cluster

The screenshot displays the MobaXterm 9.4 application window. The title bar reads "hpc.alphacruncher.net (scheideggers)". The menu bar includes Terminal, Sessions, View, X server, Tools, Games, Settings, Macros, and Help. The toolbar contains icons for Session, Servers, Tools, Games, Sessions, View, Split, MultiExec, Tunneling, Packages, Settings, Help, X server, and Exit. A "Quick connect..." field is present. On the left, a file tree shows the directory structure of the remote host, including .cache, .cm, .config, .java, .local, .matlab, .mozilla, .pki, .ssh, .subversion, .vim, intel, recycle_workshop, wmsc, .bash_history, .bash_logout, .bash_profile, .bashrc, .emacs, .kshrc, .viminfo, and .Xauthority. The main terminal window shows the following output:

```
• MobaXterm 9.4 •
(SSH client, X-server and networking tools)

> SSH session to scheideggers@hpc.alphacruncher.net
• SSH compression : ✓
• SSH-browser : ✓
• X11-forwarding : ✓ (remote display is forwarded through SSH)
• DISPLAY : ✓ (automatically set on remote server)

> For more info, ctrl+click on help or visit our website

Last login: Wed Jan 18 12:29:22 2017 from acsh.alphacruncher.net
Welcome to Bright Cluster Manager 7.3

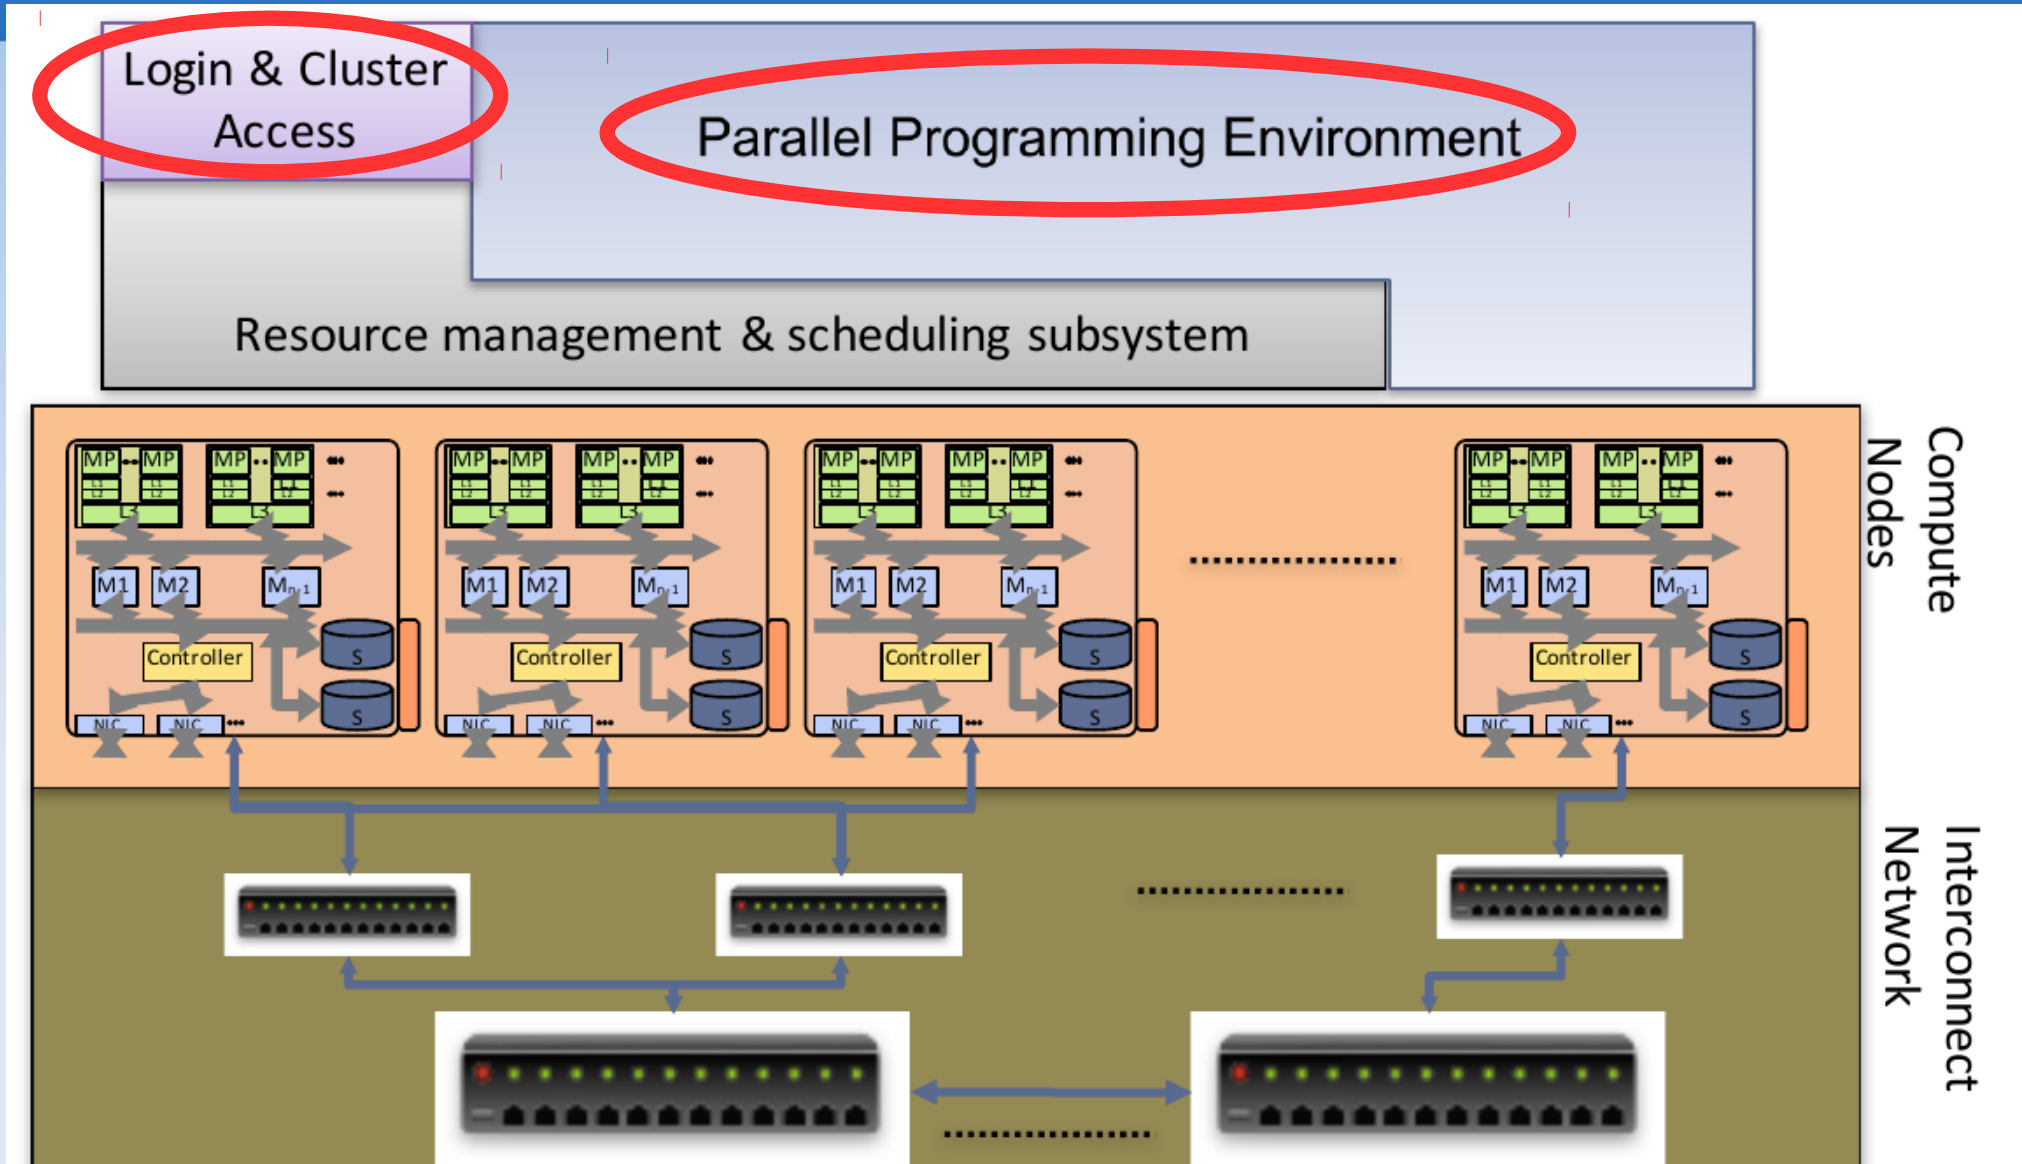
Based on CentOS Linux 7
Cluster Manager ID: #000002

Use the following commands to adjust your environment:
'module avail' - show available modules
'module add <module>' - adds a module to your environment for this session
'module initadd <module>' - configure module to be loaded at every login

[scheideggers@hpc ~]$
```

At the bottom of the window, a status bar indicates "UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: <http://mobaxterm.mobatek.net>". The system clock shows "Wednesday, January 18, 2017" and "12:48 PM 1/18/2017".

High-level architecture of a HPC system



Programming environment

Supporting diverse user community requires supporting diverse tool sets (versions of compilers, debuggers, libraries, etc)

- Check available packages

\$module avail

```
[scheideggers@hpc ~]$ module avail
----- /cm/local/modulefiles -----
cluster-tools/7.3      dot                module-info
cmd                   freeipmi/1.5.2      null
cm-scale-cluster/7.3  gcc/6.1.0          openldap
cmsh                  ipmitool/1.8.17     shared
cm-upgrade/7.3        module-git

----- /cm/shared/modulefiles -----
acml/gcc/64/5.3.1      intel/mkl/64/11.1/2013_sp1.3.174
acml/gcc/fma4/5.3.1    intel-cluster-checker/2.2.2
```

- Load a module (e.g. Matlab)

\$module load matlab/R2016b

```
[scheideggers@hpc ~]$ module load matlab/R2016b
matlab/R2015b matlab/R2016b
[scheideggers@hpc ~]$ module load matlab/R2016b
[scheideggers@hpc ~]$ module li
Currently Loaded Modulefiles:
  1) shared
  2) slurm/16.05.2
  3) openmpi/gcc/64/1.10.1
  4) python/2.7.11
  5) matlab/R2016b
[scheideggers@hpc ~]$
```

- Check what is loaded

\$module list

Note: most high-performance clusters set-up in this way !

Basic Linux commands (1)

Small-scale cheat-sheet on terminal commands

Command	Description
<code>pwd</code>	Print name of current/working directory
<code>cd [Directory]</code>	Change directory (no directory → change to home)
<code>ls [Directory]</code>	List directory contents (no directory → list current)
<code>cat FILE</code>	Concatenate files and print on the standard output
<code>mkdir DIRECTORY</code>	Make directories
<code>mkdir -p DIRECTORY</code>	Make directories, make parent directories as needed
<code>cp SOURCE... DIRECTORY</code>	Copy files and directories
<code>cp -r SOURCE... DIRECTORY</code>	Copy files and directories, copy directories recursively
<code>mv SOURCE... DIRECTORY</code>	Move (rename) files
<code>man COMMAND</code>	An interface to the on-line reference manuals

Basic Linux commands (2)

Small-scale cheat-sheet on terminal commands

Command	Description
<code>ssh -X foo@host.com</code>	OpenSSH SSH client (remote login program), access to host.com with user foo
<code>scp foo@host.com:/home/bar ./</code>	Secure copy (remote file copy program), copy file bar from /home on host.com to directory
<code>scp bar foo@host.com:/home/</code>	Secure copy (remote file copy program), copy file bar from the local host to /home on host.com
<code>git clone git@github.com:whatever folder-name</code>	The stupid content tracker, Clone a repository (whatever) into a new directory (folder-name).
<code>git checkout</code>	Checkout a branch or paths to the working tree.

Step-by-Step – an example

```
> pwd
/beegfs/swift/alphacruncher.net/USERNAME
> mkdir -p firstFolder/secondFolder
> ls
FirstFolder
> ls firstFolder
secondFolder
> cd firstFolder
> pwd
/beegfs/swift/alphacruncher.net/USERNAME/firstFolder
> ls
secondFolder
```



**KEEP
CALM
AND
GOOD
LUCK**